

A Defensible Strategy for Agentic AI Adoption

Platform selection, citizen agents and enterprise capabilities

Equip people. Connect systems. Validate the work.

Ian Purton · Bionic GPT · September 2026

A proposed strategy grounded in research, company reports and technical documentation.
Sources qualify individual claims; they do not endorse this entire strategy.

There are savings to be made.

Start with documented results, then establish your own baseline.

14%

The 2023 NBER working paper reported roughly 14% more issues resolved per hour in customer support. Gains varied substantially by worker.

Work redesign

McKinsey associates stronger AI outcomes with workflow redesign, leadership commitment and operational discipline.

Local proof

Measure time, quality and review effort in your own work. Released capacity only becomes cash savings if costs actually fall.

Treat published results as evidence of possibility, not your business case.

Evidence: one workplace study and a broader self-reported survey; neither proves universal agent savings.

1. NBER, [Generative AI at Work](#) (2023 working paper)

2. McKinsey, [The State of AI](#) (2026; respondent survey)

Make citizen agents a practical default.

For everyday knowledge work, let employees describe the work and reuse approved capabilities.

Business-owned

An employee defines the objective, source information, expected output and review rules.

Platform-enabled

Shared tools provide the models, files, integrations and execution needed to complete the task.

Real precedent

Prosus reports 60,000+ agents built across 40,000+ employees. This is company-reported adoption, not independent proof of ROI.

Begin with employee-led, reviewable tasks; increase autonomy only after validation.

Recommendation supported by a published enterprise example; citizen agent is the terminology used in this draft.

1. Prosus, *The Coming Age of AI Colleagues* (2026; company-reported)

Move from answers to completed work.

Define the outcome before choosing how much autonomy it needs.

Research

Collect approved sources and produce a referenced briefing.

Analyse

Reconcile a spreadsheet against source data and flag discrepancies.

Act

Find a customer update, propose the correct CRM change and request approval.

Use a simple workflow when the steps are fixed; use an agent when the route needs flexibility.

You need a platform and training.

Technology makes work possible. People make it useful.

Platform

Approved models, usable tools, enterprise access and controls. A platform team owns the shared foundation.

Training

Teach task definition, source checking, output review and responsible delegation using real departmental work.

Advocates

Support local champions with office hours, reusable examples and a route to report missing capabilities.

Fund enablement and operations alongside the technology.

Recommendation informed by the sources below.

1. Prosus, *The Coming Age of AI Colleagues* (2026; company-reported)

2. McKinsey, *The State of AI* (2026; respondent survey)

Control the cost of completed work.

An LLM gateway provides a common control point for model access and consumption.

See consumption

Attribute tokens, requests, latency and estimated spend to teams and tasks.
Include failed runs and retries.

Set limits

Use quotas, rate limits and alerts. Keep provider credentials out of agent instructions and user prompts.

Choose models

Test lower-cost models on representative tasks. Change defaults only when quality and reliability remain acceptable.

Track cost per accepted result, not just the price per token.

Microsoft documents gateway controls; the selection and measurement approach is our recommendation.

1. Microsoft Learn, [AI gateway capabilities](#)

2. Prosus, [The Coming Age of AI Colleagues](#) (2026; company-reported)

Budget beyond model tokens.

The platform cost includes the people and infrastructure needed to operate it.

Recurring cost

Licences or inference, compute, storage, gateways, monitoring, support and platform operations.

Change cost

Integrations, training, evaluation, upgrades and migration. Include internal engineering time.

Useful denominator

Monthly operating cost divided by accepted completed tasks. Track quality and task mix alongside it.

Self-hosted does not automatically mean cheaper; compare fully loaded costs.

Proposed costing method. No vendor price comparison or savings forecast is implied.

1. McKinsey, [The State of AI \(2026; respondent survey\)](#)

Build a reusable integration catalogue.

Connect capabilities once, then make them available to multiple tasks.

Find the systems

Inventory the systems of record, API owners, available operations and access constraints.

Describe operations

Provide schemas, clear descriptions, examples and error behaviour. Prefer existing API contracts where practical.

Expose selectively

Offer approved operations by role and task. Test discovery and correct use, not just connectivity.

A connector list is not enough: the required operation must work with the right permissions.

API management and tool-interface guidance support this proposed integration approach.

1. [Microsoft Learn, API Management overview](#)

2. [Anthropic, Building Effective Agents \(architecture guidance\)](#)

Put API governance in the execution path.

An API gateway can enforce controls; governance defines what those controls should be.

Governance

Name owners. Define permitted use, write approvals, API versions, quotas and change responsibilities.

Gateway

Enforce credentials, routing, rate limits and telemetry. Reuse an existing enterprise gateway where it fits.

System authority

Preserve record-level permissions in the destination system. Gateway authentication alone is insufficient.

The agent can request an action; policy and the target system decide whether it is allowed.

Microsoft documents gateway functions. The governance and permission recommendations extend that foundation.

1. [Microsoft Learn, API Management overview](#)

Give agents tools to produce usable work.

Files, execution and artifact generation determine which tasks the platform can complete.

Work with files

Read and write spreadsheets, documents and PDFs. Preserve formulas, provenance and editable outputs.

Execute safely

Provide an isolated runtime with limits on time, resources, network access and secrets.

Match the workload

Test a lightweight runtime for routine tasks; add a fuller sandbox when required libraries or tools need it.

Choose the execution environment from tested task requirements.

Proposed capability requirements; the sources do not establish a universal sandbox design.

1. Anthropic, [Building Effective Agents](#) (architecture guidance)

2. NIST, [Generative AI Profile, AI 600-1](#) (2024)

Keep identity and data access intact.

Agent access should follow the permissions and data boundaries of the organisation.

Identity

Use enterprise sign-in and scoped credentials. Define whose authority a scheduled job uses.

Knowledge

Respect document permissions in search and retrieval. Show source references and recognise stale information.

Data boundaries

Approve where prompts, files, logs and model traffic can go. Set retention and deletion policies.

Connecting data is only useful when access, provenance and handling are trustworthy.

Recommended controls, informed by NIST risk management and API management guidance.

1. NIST, [Generative AI Profile, AI 600-1 \(2024\)](#)

2. Microsoft Learn, [API Management overview](#)

Make recurring work operable.

Scheduling turns a useful task into a service that someone must own.

Run controls

Schedules or triggers, bounded retries, timeouts and protection against duplicate side effects.

Lifecycle

An owner, versioned instructions, test cases, approved tools and a clear pause or retirement mechanism.

Human review

Approval before consequential actions; failure notifications with enough context for someone to intervene.

Every recurring agent needs an owner and a way to stop it.

Operational requirements proposed for this strategy, rather than a feature list proven by the sources.

1. NIST, [Generative AI Profile, AI 600-1 \(2024\)](#)

2. Anthropic, [Building Effective Agents \(architecture guidance\)](#)

Evaluate and observe the whole task.

A successful model response is only one part of a successful business result.

Before release

Test representative tasks, incorrect inputs, permission failures and unexpected tool responses.

During operation

Record outcome, latency, model version, tool errors, review effort and cost. Redact sensitive logging.

After changes

Repeat the relevant acceptance set when models, prompts, tools or permissions change.

Require evidence that the platform completes the work reliably.

Recommendation informed by the sources below.

1. NIST, Generative AI Profile, AI 600-1 (2024)
2. Anthropic, Building Effective Agents (architecture guidance)

Choose the deployment approach.

Start from constraints, then compare the operating model.

Managed SaaS

Consider when approved data handling and available integrations meet the task requirements. Validate access to private systems.

Self-hosted platform

Consider when customer-controlled infrastructure, model choice or internal integration needs justify operating the platform.

Custom development

Consider when validated requirements remain unmet or a differentiated workflow warrants dedicated engineering.

Use the least complex approach that satisfies your acceptance criteria.

Decision guidance. Bionic is an example of the self-hosted category, based on its own description.

1. Anthropic, [Building Effective Agents](#) (architecture guidance)

2. Bionic GPT, [product overview](#) (vendor description)

Preserve the ability to change.

Portability must be tested at the task level.

Models

Keep a repeatable acceptance set and exportable model settings. A compatible API does not guarantee equivalent behaviour.

Tools and knowledge

Retain API contracts, source documents, permission mappings and integration documentation.

Agent definitions

Export instructions, schedules and evaluation cases. Estimate the effort to recreate platform-specific features.

Ask for a migration demonstration, not just a claim of model independence.

Proposed procurement criteria; portability depends on the implementation and contract.

1. [Microsoft Learn, API Management overview](#)

2. [Anthropic, Building Effective Agents \(architecture guidance\)](#)

Compare against explicit requirements.

Complete this worksheet with demonstrations and contractual evidence.

Criterion	Managed SaaS	Bionic GPT	Custom platform
Deployment	Confirm approved hosting	On-prem / private cloud ¹	Your architecture
Models	Confirm available options	Hosted / private / local ¹	Your integrations
Business systems	Demonstrate required actions	OpenAPI integration ¹ ; test actions	Implement and operate
Files + execution	Run acceptance tasks	Run acceptance tasks	Build and test
Identity + controls	Verify end-to-end	Verify end-to-end	Design and verify
Scheduling + recovery	Demonstrate failure handling	Demonstrate failure handling	Build and operate
Model + API gateways	Inspect native / external fit	Inspect native / external fit	Choose and integrate
Portability	Test exports and migration	Test exports and migration	Document and test
Total cost + support	Contract + internal operations	Infrastructure + support + operations	Engineering + operations

Mark each requirement: meets / needs extension / fails, with evidence and an owner.

¹ Bionic product claims, not an independent feature assessment. Other cells are due-diligence prompts.
Author affiliation: Bionic GPT. Compare named candidates within each category before deciding.
[1. Bionic GPT product overview](#)

VALIDATION

Test the platform on your work.

Use a small acceptance set spanning the capabilities you need.

Documents + data

Create a briefing with verifiable citations; reconcile a workbook and preserve formulas.

Enterprise actions

Read and update a test record with correct identity, permissions and approval behaviour.

Recurring execution

Run a scheduled task through failures, retries and notification. Check for duplicate writes.

Set task-specific quality, reliability, cost and review thresholds before comparing options.

Illustrative acceptance tasks. Passing them supports only the capabilities and conditions actually tested.

1. NIST, [Generative AI Profile, AI 600-1 \(2024\)](#)

ENABLEMENT

Train people to delegate and review.

Teach a repeatable method using their own work.

Describe

State the objective, inputs, output format, constraints and how completion will be judged.

Verify

Check facts, calculations, sources and proposed actions. Learn when to request specialist help.

Reuse

Turn validated instructions into shared examples; nominate an owner before scheduling or broad distribution.

Measure whether training produces independently completed, reviewed tasks.

Proposed training curriculum, informed by enterprise adoption and workflow-redesign evidence.

1. Prosus, *The Coming Age of AI Colleagues* (2026; company-reported)

2. McKinsey, *The State of AI* (2026; respondent survey)

Expand through evidence and shared capabilities.

Use decision gates rather than assuming every experiment should become production.

Establish

Name a sponsor and platform owner.
Baseline tasks, approve data boundaries and connect initial systems.

Validate

Train an initial group. Test useful tasks and address capability gaps. Proceed when quality and controls meet agreed thresholds.

Expand

Share validated methods, add reusable integrations and support departmental ownership. Retire unsuccessful agents.

Scale the capabilities that enable useful work, then expand the work they support.

Proposed rollout sequence. Scope and timing should follow organisational readiness.

1. Prosus, *The Coming Age of AI Colleagues* (2026; company-reported)

2. NIST, *Generative AI Profile, AI 600-1* (2024)

OUTCOMES

Measure accepted work and realised value.

Usage shows adoption. Outcomes show whether the investment is worthwhile.

Delivery

Accepted completion rate, task volume, quality checks, human review and recovery effort.

Economics

Total operating cost; net time saved after review; cost per accepted task; evidenced revenue or cash effects.

Adoption

Active repeat users, independently completed tasks and reuse of approved integrations and instructions.

Keep capacity released, cash savings and revenue gains as separate measures.

Proposed scorecard. Self-reported time savings need verification before financial attribution.

1. McKinsey, *The State of AI* (2026; respondent survey)

2. NIST, *Generative AI Profile, AI 600-1* (2024)

Keep the argument and evidence distinct.

How to read the sources

NBER supports a bounded productivity claim, not a universal saving.
Prosus describes its own deployment and observations. McKinsey reports survey evidence and associations. Microsoft documents implementation capabilities. NIST provides risk-management guidance. Anthropic offers architecture guidance. Bionic describes its own product.

Before publishing

Rewrite the recommendation in your own voice. Choose the target audience. Add organisation-specific thresholds and examples. Complete the matrix for named candidates. Recheck changing product features and sources. Keep the affiliation and evidence qualifications visible.

The strategy is an argument to test, not a promise of a particular financial return.